

Dynamic Memory Allocation

15-213: Introduction to Computer Systems
Recitation 11: Monday, Nov 3, 2014

SHAILIN DESAI
SECTION L

Today

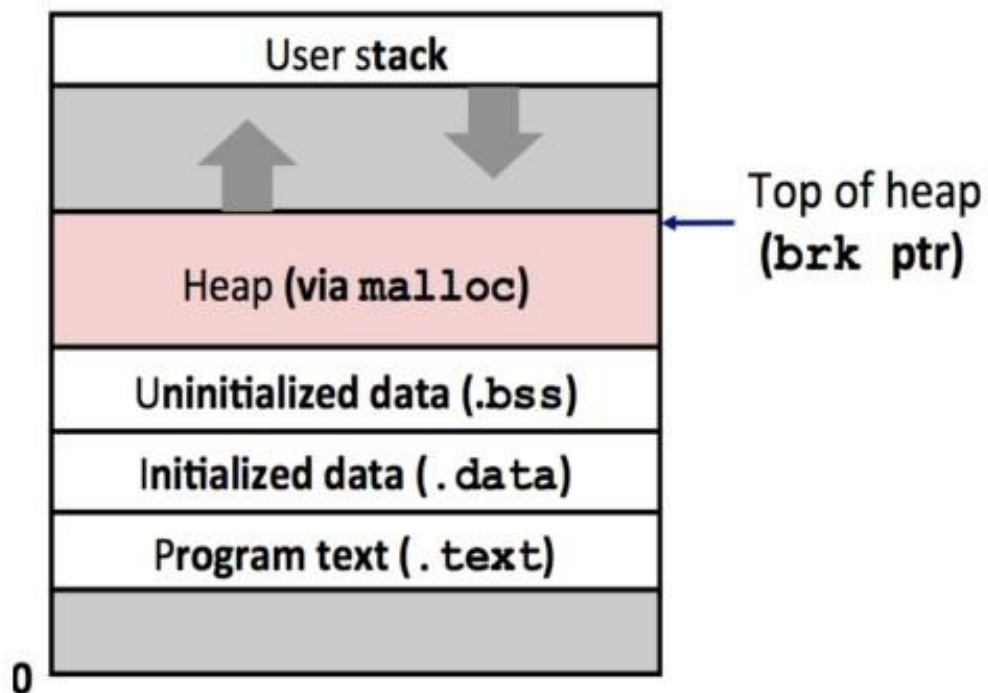
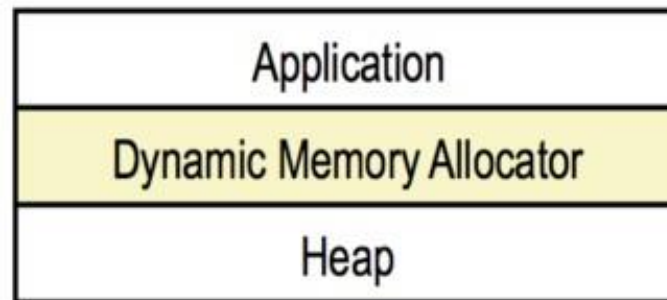
- Lecture Review
- Macros and Inline Functions
- Malloc Lab
- Heap Checker

Today

- *Lecture Review*
- Macros and Inline Functions
- Malloc Lab
- Heap Checker

Dynamic Memory Allocation

- Programmers use *dynamic memory allocators* (such as `malloc`) to acquire VM at run time.
- Dynamic memory allocators manage an area of process virtual memory known as the *heap*.



Dynamic Memory Allocation

```
p1 = malloc(4)
```




```
p2 = malloc(5)
```




```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(2)
```



How do we know where to put the next block?

Keeping Track of Free Blocks

- Method 1: **Implicit list** using length—links all blocks



- Method 2: **Explicit list** among the free blocks using pointers



- Method 3: **Segregated free list**
 - *Different free lists for free blocks of different size classes*

Method 1: Implicit List


- **For each block, we need both size and allocation status**

Could store this information in two words: wasteful!

- **Standard trick**

If blocks are aligned, some low-order address bits are always 0
Instead of storing an always-0 bit, use it as a allocated/free flag

*Format of
allocated and
free blocks*



a = 1: Allocated block
a = 0: Free block

Size: block size

**Payload: application data
(allocated blocks only)**

Method 2: Explicit List

➤ Maintain list(s) of *free* blocks instead of *all* blocks


The “next” free block could be anywhere

So we need to store forward/back pointers, not just sizes


Still need boundary tags for coalescing

➤ *Luckily we track only free blocks, so we can use payload area*

Allocated (as before)



Free




Method 2: Explicit Free Lists

➤ Logically...




➤ But physically...



Method 3: Segregated List

- Each *size class* of blocks has its own free list



- Small sized blocks: more lists for separate classes
- Larger sizes: one class for each two-power size

Finding a Free Block


➤ First fit:

Search list from beginning, choose first free block that fits

Can take linear time in total number of blocks (allocated and free)

➤ In practice it can cause “splinters” at beginning of list



Many small free blocks left at beginning



Finding a Free Block

➤ Next fit:

- Like first fit, but search list starting where previous search finished
- Should often be faster than first fit: avoids re-scanning unhelpful blocks
- Some research suggests that fragmentation is worse

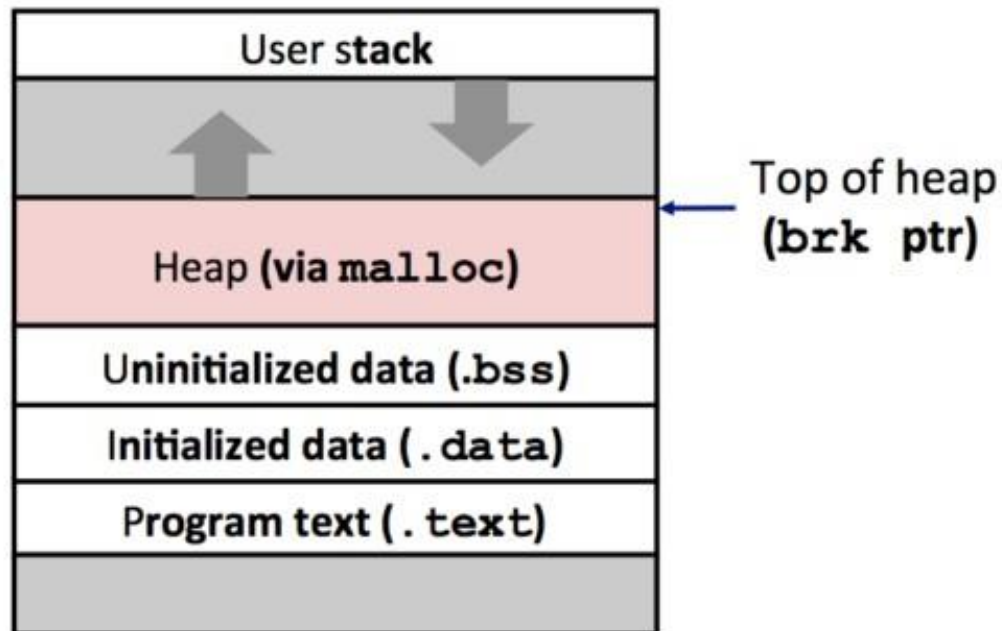


Finding a Free Block

- **Best fit:**
 - Search the list, choose the best free block: fits, with fewest bytes left over
 - Keeps fragments small: usually improves memory utilization
 - Will typically run slower than first fit
- **If the block we find is larger than we need, split it**

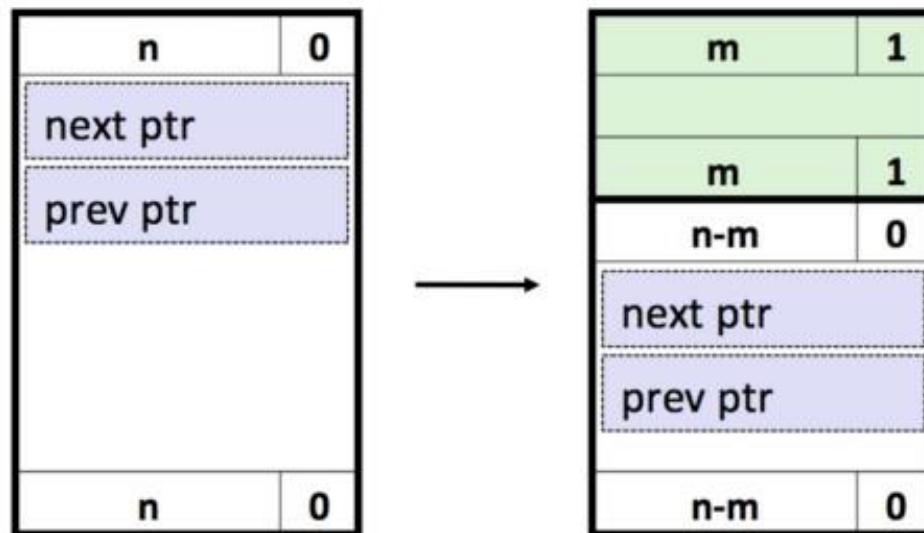
Finding a Free Block

- **What happens if we can't find a block?**
 - Need to extend the heap
 - Use the `brk()` or `sbrk()` system calls
 - In `mallocLab`, use `mem_sbrk()`
 - `sbrk(requested space)` allocates space and returns pointer to start of new space
 - `sbrk(0)` returns pointer to top of current heap
 - Use what you need, add the rest as a whole free block



Splitting a Block

- **What happens if the block we have is too big?**
 - Split between portion we need and the leftover free space
 - For implicit lists: correct the block size
 - For explicit lists: correct the previous and next pointers
 - For segregated lists:
 - determine correct size list
 - Insert with insertion policy (more on this later)



Freeing Blocks


➤ Simplest implementation:

- Need only clear the “allocated” flag

`void free_block(ptr p) { *p = *p & -2 }`

- But can lead to external fragmentation:

- There is enough free space, but the allocator can't find it



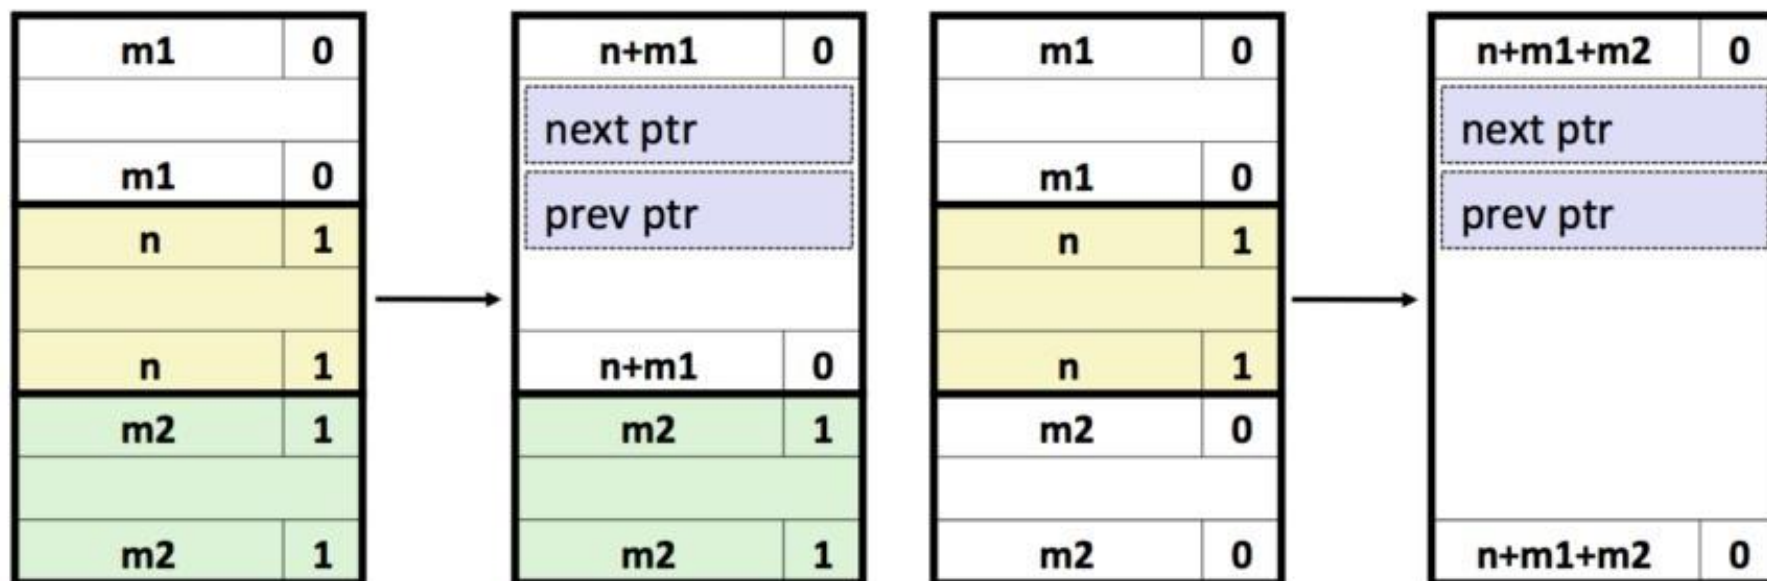
Freeing Blocks

- **Need to combine blocks nearby in memory (coalescing)**
- **For implicit lists:**
 - Simply look backwards and forwards using block sizes
- **For explicit lists:**
 - Look backwards/forwards using block sizes, not next/prev pointers
- **For segregated lists:**
 - use the size of new block to determine proper list
 - Insert back into list based on insertion policy (LIFO, FIFO)



Freeing Blocks

- **Graphical depiction (both implicit & explicit):**
- (these are physical mappings)



Insertion Policy

➤ **Where in the free list do you put a newly freed block?**

➤ **LIFO (last-in-first-out) policy**

- Insert freed block at the beginning of the free list
- **Pro:** simple and constant time
- **Con:** studies suggest fragmentation is worse than address ordered

➤ **Address-ordered policy**

- Insert freed blocks so that free list blocks are always in address order:
 - $addr(prev) < addr(curr) < addr(next)$
- **Con:** requires search
- **Pro:** fragmentation is lower than LIFO

Today

- Lecture Review
- *Macros and Inline Functions*
- Malloc Lab
- Heap Checker

Macros

- C Preprocessor looks at macros in the preprocessing step of compilation
- Use `#define` to avoid magic numbers:
 - `#define TRIALS 100`
- Function like macros - short and heavily used code snippets
 - `#define GET_BYTE_ONE(x) ((x) & 0xff)`
 - `#define GET_BYTE_TWO(x) (((x) >> 8) & 0xff)`
- Inline functions
 - Ask the compiler to insert the complete body of the function in every place that the function is called (simply replacing code)
 - `inline int fun(int a, int b)`
 - Requests compiler to insert assembly of fun wherever a call to fun is made
- Both are useful for mallocab

Assert()

- `assert(expr)`
 - If `expr` is false, the calling process is terminated
 - If `expr` is true, it does nothing
- May be turned off at compile time with option `-DNDEBUG`
- As always, “Man is your friend.”
- For style points: you **MUST** use asserts in your code

Debugging

➤ Using printf, assert, etc only in debug mode:

- `#define DEBUG`
- `#ifdef DEBUG`
 - `# define dbg_printf(...) printf(__VA_ARGS__)`
 - `# define dbg_assert(...) assert(__VA_ARGS__)`
 - `# define dbg(...)`
- `#else`
 - `# define dbg_printf(...)`
 - `# define dbg_assert(...)`
 - `# define dbg(...)`
- `#endif`

Today

- Lecture Review
- Macros and Inline Functions
- *Malloc Lab*
- Heap Checker

Malloclab

➤ **You need to implement the following functions:**

- `int mm_init(void);`
- `void *malloc(size_t size);`
- `void free(void *ptr);`
- `Void *realloc(void *ptr, size_t size);`
- `void *calloc (size_t n, size_t size);`
- `void mm_checkheap(int verbose);`

➤ **Scored on space efficiency and throughput**

➤ **Cannot call system memory functions**

➤ **Use helper functions (as static/inline functions)**

➤ **May want to consider practicing version control**

Malloclab

➤ Inline

- Essentially copies function code into location of each function call
- Avoids overhead of stack discipline/function call (once assembled)
- Can often be used in place of macros
- Strong type checking and input variable handling, unlike macros.

➤ Static

- Resides in a single place in memory
- Limits scope of function to the current translations unit (file)
- Should use this for helper functions only called *locally*
- Avoids polluting namespace.

➤ **static inline**

- Not surprisingly, can be used together

Today

- Lecture Review
- Macros and Inline Functions
- Malloc Lab
- *Heap Checker*

Heap Checker

- **Int mm_checkheap(int verbose) is critical for debugging**
 - Write this early
 - update it when you change your free list implementation
 - It should ensure that you haven't lost control of any part of heap memory (everything should either be allocated or listed)
- Look over lecture notes on garbage collection (particularly mark & sweep).
- This function is meant to be correct, not efficient.

Heap Checker

- Once you've settled on a design, write the heap checker that checks all the invariants of the particular design
- The checking should be detailed enough that the heap check passes if and only if the heap is truly well-formed
- Call the heap checker before/after the major operations whenever the heap should be well-formed
- Define macros to enable/disable it conveniently

- e.g.

```
#ifdef DEBUG
#define CHECKHEAP(verbose) printf("%s\n", __func__); mm_checkheap(verbose);
#endif
```

Heap Checker



- The `mm_checkheap` function takes a single integer argument that you can use any way you want.
- One very useful technique is to use this argument to pass in the line number of the call site:
 - `mm_checkheap(__LINE__);`
- If `mm_checkheap` detects a problem with the heap, it can print the line number where `mm_checkheap` was called, which allows you to call `mm_checkheap` at numerous places in your code while you are debugging.

Invariants (non-exhaustive)

- Block level:
 - Header and footer match
 - Payload area is aligned
- List level:
 - Next/prev pointers in consecutive free blocks are consistent
 - Free list contains no allocated blocks
 - All free blocks are in the free list
 - No contiguous free blocks in memory (unless you defer coalescing)
 - No cycles in the list (unless you use circular lists)
 - Segregated list contains only blocks that belong to the size class
- Heap level:
 - Prologue/Epilogue blocks are at specific locations (e.g. heap boundaries) and have special size/alloc fields
 - All blocks stay in between the heap boundaries
- And your own invariants (e.g. address order)

Hare and Tortoise Algorithm

- Detects cycles in linked lists
- Set two pointers “hare” and “tortoise” to the beginning of the list
- During each iteration, move the hare pointer forward two nodes and move the tortoise forward one node. If they are pointing to the same node after this, the list has a cycle.
- If the tortoise reaches the end of the list, there are no cycles.



Asking for help

- It can be hard for the TAs to debug your allocator, because this is a more open-ended lab
- Before asking for help, ask yourself some questions:
 - What part of which trace file triggers the error?
 - Around the point of the error, what sequence of events do you expect?
 - What part of the sequence already happened?
- If you can't answer, it's a good idea to gather more information...
 - How can you measure which step worked OK?
 - printf, breakpoints, watchpoints...

Debugging

➤ Valgrind!

- Powerful debugging and analysis technique
- Rewrites text section of executable object file
- Can detect all errors as debugging **malloc**
- Can also check each individual reference at runtime
 - Bad pointers
 - Overwriting
 - Referencing outside of allocated block

➤ GDB

- You know how to use this (hopefully)

Beyond Debugging: Error prevention

- It is hard to write code that are completely correct the first time, but certain practices can make your code less error-prone
- Plan what each function does before writing code
 - Draw pictures when linked list is involved
 - Consider edge cases when the block is at start/end of list
- Document your code as you write it

Questions?

➤ Good luck!:D