

# **Machine-Level Programming I: Basics**

15-213/18-213/14-513/15-513: Introduction to Computer Systems 5<sup>th</sup> Lecture, September 11, 2018

# **Today: Machine Programming I: Basics**

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- C, assembly, machine code

## **Intel x86 Processors**

## Dominate laptop/desktop/server market

### Evolutionary design

- Backwards compatible up until 8086, introduced in 1978
- Added more features as time goes on

## Complex instruction set computer (CISC)

- Many different instructions with many different formats
  - But, only small subset encountered with Linux programs
- Hard to match performance of Reduced Instruction Set Computers (RISC)
- But, Intel has done just that!
  - In terms of speed. Less so for low power.

## Intel x86 Evolution: Milestones

Name Date Transistors MHz

■ 8086 1978 29K 5-10

First 16-bit Intel processor. Basis for IBM PC & DOS

1MB address space

■ 386 1985 275K 16-33

First 32 bit Intel processor, referred to as IA32

Added "flat addressing", capable of running Unix

■ Pentium 4E 2004 125M 2800-3800

First 64-bit Intel x86 processor, referred to as x86-64

■ Core 2 2006 291M 1060-3333

First multi-core Intel processor

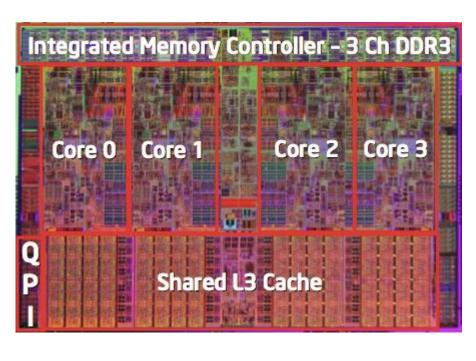
■ Core i7 2008 731M 1600-4400

Four cores (our shark machines)

## Intel x86 Processors, cont.

#### Machine Evolution

<b>386</b>	1985	0.3M
Pentium	1993	3.1M
Pentium/MMX	1997	4.5M
PentiumPro	1995	6.5M
Pentium III	1999	8.2M
Pentium 4	2000	42M
Core 2 Duo	2006	291M
Core i7	2008	731M



#### Added Features

Core i7 Skylake

Instructions to support multimedia operations

2015

Instructions to enable more efficient conditional operations

1.9B

- Transition from 32 bits to 64 bits
- More cores

## Intel x86 Processors, cont.

#### ■ Past Generations Process technology

■ 1<sup>st</sup> Pentium Pro 1995 600 nm

■ 1<sup>st</sup> Pentium III 1999 250 nm

■ 1<sup>st</sup> Pentium 4 2000 180 nm

■ 1<sup>st</sup> Core 2 Duo 2006 65 nm

## ■ Recent & Upcoming Generations

1. Nehalem 2008 45 nm

2. Sandy Bridge 2011 32 nm

3. Ivy Bridge 2012 22 nm

4. Haswell 2013 22 nm

5. Broadwell 2014 14 nm

6. Skylake 2015 14 nm

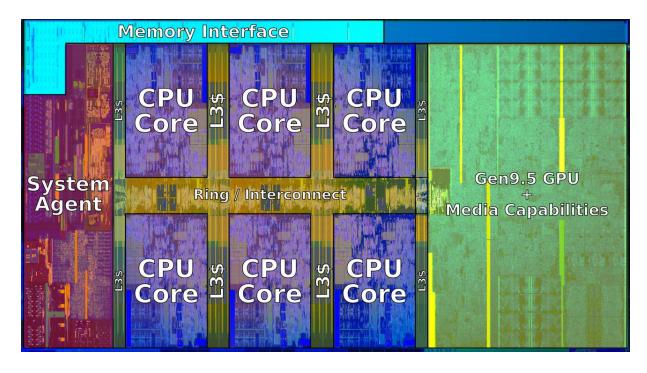
7. Kaby Lake 2016 14 nm

8. Coffee Lake 2017 14 nm

Cannon Lake 2019? 10 nm

Process technology dimension= width of narrowest wires(10 nm ≈ 100 atoms wide)

## 2018 State of the Art: Coffee Lake



#### ■ Mobile Model: Core i7

- 2.2-3.2 GHz
- **45 W**

## Desktop Model: Core i7

- Integrated graphics
- 2.4-4.0 GHz
- **35-95 W**

#### ■ Server Model: Xeon E

- Integrated graphics
- Multi-socket enabled
- 3.3-3.8 GHz
- **80-95 W**

# x86 Clones: Advanced Micro Devices (AMD)

### Historically

- AMD has followed just behind Intel
- A little bit slower, a lot cheaper

#### Then

- Recruited top circuit designers from Digital Equipment Corp. and other downward trending companies
- Built Opteron: tough competitor to Pentium 4
- Developed x86-64, their own extension to 64 bits

#### Recent Years

- Intel got its act together
  - Leads the world in semiconductor technology
- AMD has fallen behind
  - Relies on external semiconductor manufacturer

# Intel's 64-Bit History

- 2001: Intel Attempts Radical Shift from IA32 to IA64
  - Totally different architecture (Itanium)
  - Executes IA32 code only as legacy
  - Performance disappointing
- 2003: AMD Steps in with Evolutionary Solution
  - x86-64 (now called "AMD64")
- Intel Felt Obligated to Focus on IA64
  - Hard to admit mistake or that AMD is better
- 2004: Intel Announces EM64T extension to IA32
  - Extended Memory 64-bit Technology
  - Almost identical to x86-64!
- All but low-end x86 processors support x86-64
  - But, lots of code still runs in 32-bit mode

## **Our Coverage**

#### ■ IA32

- The traditional x86
- For 15/18-213: RIP, Summer 2015

#### ■ x86-64

- The standard
- shark> gcc hello.c
- shark> gcc -m64 hello.c

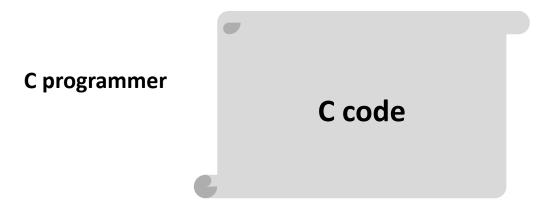
#### Presentation

- Book covers x86-64
- Web aside on IA32
- We will only cover x86-64

# **Today: Machine Programming I: Basics**

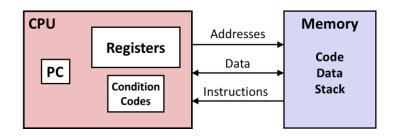
- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- C, assembly, machine code

## **Levels of Abstraction**



Nice clean layers, but beware...

#### **Assembly programmer**





**Computer Designer** 

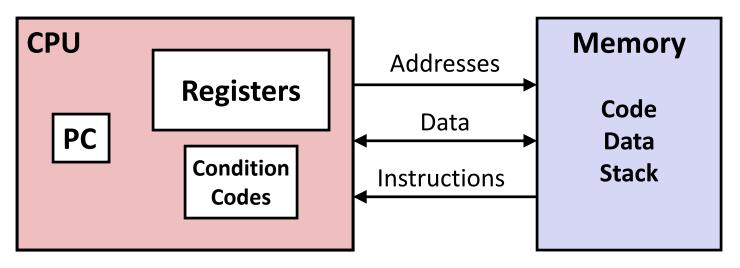
Caches, clock freq, layout, ...

Of course, you know that: It's why you are taking this course.

## **Definitions**

- Architecture: (also ISA: instruction set architecture) The parts of a processor design that one needs to understand for writing correct machine/assembly code
  - Examples: instruction set specification, registers
  - Machine Code: The byte-level programs that a processor executes
  - Assembly Code: A text representation of machine code
- Microarchitecture: Implementation of the architecture
  - Examples: cache sizes and core frequency
- Example ISAs:
  - Intel: x86, IA32, Itanium, x86-64
  - ARM: Used in almost all mobile phones
  - RISC V: New open-source ISA

## **Assembly/Machine Code View**



## **Programmer-Visible State**

- PC: Program counter
  - Address of next instruction
  - Called "RIP" (x86-64)
- Register file
  - Heavily used program data
- Condition codes
  - Store status information about most recent arithmetic or logical operation
  - Used for conditional branching

#### Memory

- Byte addressable array
- Code and user data
- Stack to support procedures

Bryant a

# **Assembly Characteristics: Data Types**

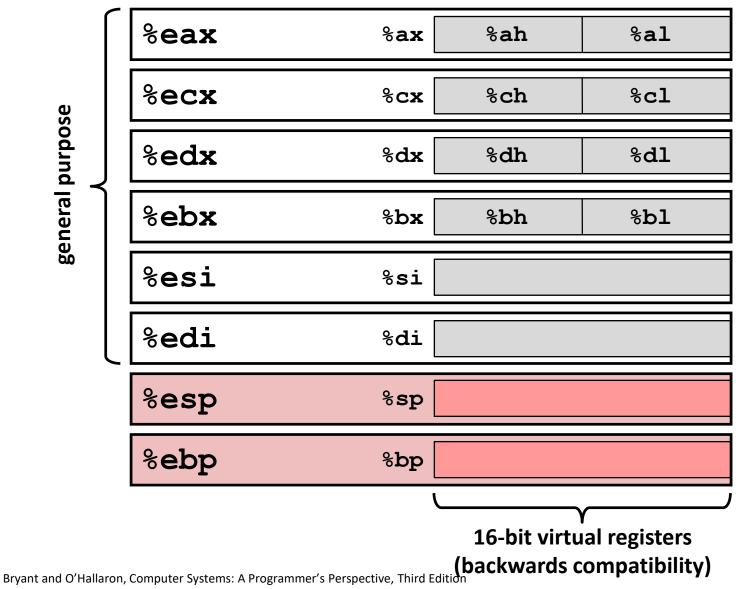
- "Integer" data of 1, 2, 4, or 8 bytes
  - Data values
  - Addresses (untyped pointers)
- **■** Floating point data of 4, 8, or 10 bytes
- (SIMD vector data types of 8, 16, 32 or 64 bytes)
- Code: Byte sequences encoding series of instructions
- No aggregate types such as arrays or structures
  - Just contiguously allocated bytes in memory

## x86-64 Integer Registers

%rax	%eax	%r8	%r8d
%rbx	%ebx	%r9	%r9d
%rcx	%ecx	%r10	%r10d
%rdx	%edx	%r11	%r11d
%rsi	%esi	%r12	%r12d
%rdi	%edi	%r13	%r13d
%rsp	%esp	%r14	%r14d
%rbp	%ebp	%r15	%r15d

- Can reference low-order 4 bytes (also low-order 1 & 2 bytes)
- Not part of memory (or cache)

# Some History: IA32 Registers



#### Origin (mostly obsolete)

accumulate

counter

data

base

source index

destination index

stack pointer base pointer

## **Assembly Characteristics: Operations**

- Transfer data between memory and register
  - Load data from memory into register
  - Store register data into memory
- Perform arithmetic function on register or memory data
- Transfer control
  - Unconditional jumps to/from procedures
  - Conditional branches
  - Indirect branches

## **Moving Data**

- Moving Data
  movq Jource, Dest
- Operand Types
  - Immediate: Constant integer data
    - Example: \$0x400, \$-533
    - Like C constant, but prefixed with `\$'
    - Encoded with 1, 2, or 4 bytes
  - Register: One of 16 integer registers
    - Example: %rax, %r13
    - But %rsp reserved for special use
    - Qthers have special uses for particular instructions
  - Memory 8 consecutive bytes of memory at address given by register
    - Simplest example: (%rax)
    - Various other "addressing modes"

%rax
%rcx
%rdx
%rbx
%rsi
%rdi
%rsp

%rN

%rbp

Warning: Intel docs use mov *Dest, Source* 

## movq Operand Combinations

Cannot do memory-memory transfer with a single instruction

# **Simple Memory Addressing Modes**

- Normal (R) Mem[Reg[R]]
  - Register R specifies memory address
  - Aha! Pointer dereferencing in C

```
movq (%rcx),%rax
```

- Displacement D(R) Mem[Reg[R]+D]
  - Register R specifies start of memory region
  - Constant displacement D specifies offset

```
movq 8(%rbp),%rdx
```

# **Example of Simple Addressing Modes**

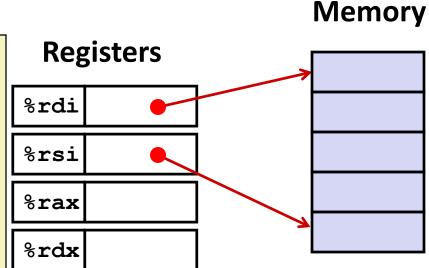
```
void
whatAmI(<type> a, <type> b)
{
    ????
}
    movq (%rdi), %rax
    movq (%rsi), %rdx
    movq %rdx, (%rdi)
    movq %rax, (%rsi)
    ret

%rdi
```

# **Example of Simple Addressing Modes**

```
void swap
   (long *xp, long *yp)
{
   long t0 = *xp;
   long t1 = *yp;
   *xp = t1;
   *yp = t0;
}
```

# void swap (long \*xp, long \*yp) { long t0 = \*xp; long t1 = \*yp; \*xp = t1; \*yp = t0; }

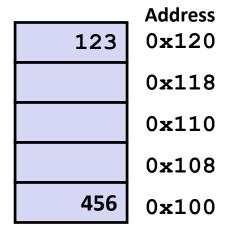


Register	Value
%rdi	хр
%rsi	ур
%rax	t0
%rdx	t1

## Registers

%rdi	0x120
%rsi	0x100
%rax	
%rdx	

#### Memory



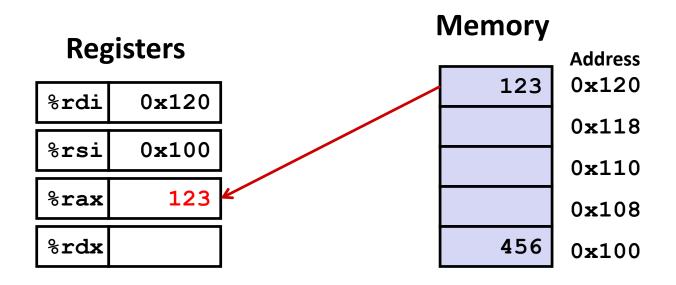
```
movq (%rdi), %rax # t0 = *xp

movq (%rsi), %rdx # t1 = *yp

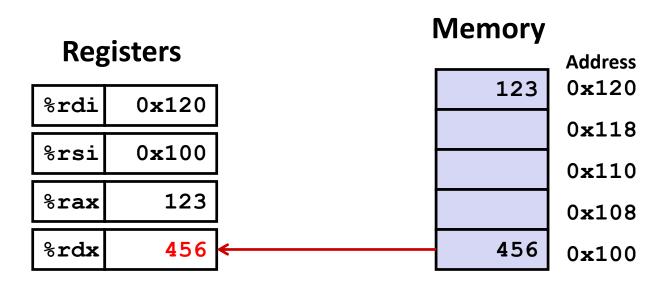
movq %rdx, (%rdi) # *xp = t1

movq %rax, (%rsi) # *yp = t0

ret
```



```
movq (%rdi), %rax # t0 = *xp
movq (%rsi), %rdx # t1 = *yp
movq %rdx, (%rdi) # *xp = t1
movq %rax, (%rsi) # *yp = t0
ret
```



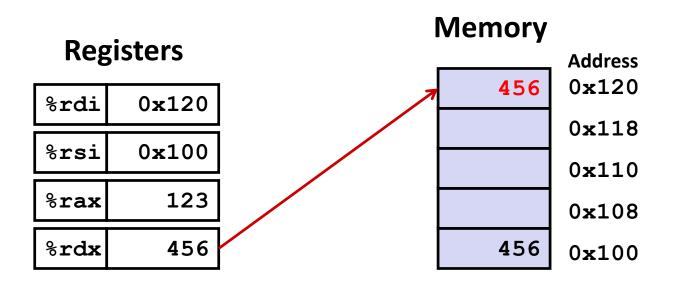
```
movq (%rdi), %rax # t0 = *xp

movq (%rsi), %rdx # t1 = *yp

movq %rdx, (%rdi) # *xp = t1
```

movq %rax, (%rsi) # \*yp = t0

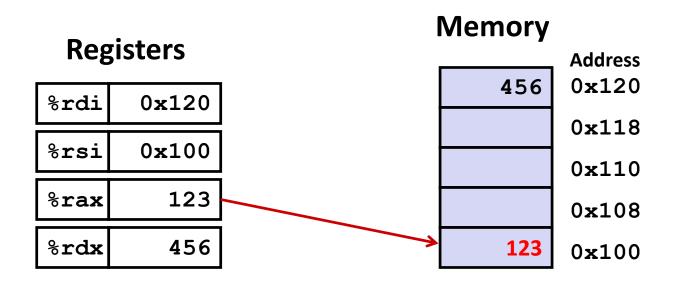
ret



```
swap:
          (\$rdi), \$rax # t0 = *xp
  movq
       (%rsi), %rdx # t1 = *yp
  movq
       %rdx, (%rdi) # *xp = t1
  movq
          %rax, (%rsi) # *yp = t0
```

ret

movq



```
movq (%rdi), %rax # t0 = *xp
movq (%rsi), %rdx # t1 = *yp
```

movq %rdx, (%rdi) # \*xp = t1
movq %rax, (%rsi) # \*yp = t0

ret

# **Simple Memory Addressing Modes**

- Normal (R) Mem[Reg[R]]
  - Register R specifies memory address
  - Aha! Pointer dereferencing in C

```
movq (%rcx),%rax
```

- Displacement D(R) Mem[Reg[R]+D]
  - Register R specifies start of memory region
  - Constant displacement D specifies offset

```
movq 8(%rbp),%rdx
```

# **Complete Memory Addressing Modes**

#### Most General Form

D(Rb,Ri,S) Mem[Reg[Rb]+S\*Reg[Ri]+D]

D: Constant "displacement" 1, 2, or 4 bytes

Rb: Base register: Any of 16 integer registers

Ri: Index register: Any, except for %rsp

S: Scale: 1, 2, 4, or 8 (why these numbers?)

#### Special Cases

(Rb,Ri) Mem[Reg[Rb]+Reg[Ri]]

D(Rb,Ri) Mem[Reg[Rb]+Reg[Ri]+D]

(Rb,Ri,S) Mem[Reg[Rb]+S\*Reg[Ri]]

## **Address Computation Examples**

%rdx	0xf000
%rcx	0x0100

D(Rb,Ri,S)	Mem[Reg[Rb]+S*Reg[Ri]+ D]
------------	---------------------------

- D: Constant "displacement" 1, 2, or 4 bytes
   Rb: Base register: Any of 16 integer registers
   Ri: Index register: Any, except for %rsp

- S: Scale: 1, 2, 4, or 8 (why these numbers?)

Expression	Address Computation	Address
0x8(%rdx)		
(%rdx,%rcx)		
(%rdx,%rcx,4)		
0x80(,%rdx,2)		

# **Address Computation Examples**

%rdx	0xf000
%rcx	0x0100

Expression	Address Computation	Address
0x8(%rdx)	0xf000 + 0x8	0xf008
(%rdx,%rcx)	0xf000 + 0x100	0xf100
(%rdx,%rcx,4)	0xf000 + 4*0x100	0xf400
0x80(,%rdx,2)	2*0xf000 + 0x80	0x1e080

# **Today: Machine Programming I: Basics**

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- C, assembly, machine code

## **Address Computation Instruction**

### ■ leaq Src, Dst

- Src is address mode expression
- Set Dst to address denoted by expression

#### Uses

- Computing addresses without a memory reference
  - E.g., translation of p = &x[i];
- Computing arithmetic expressions of the form x + k\*y
  - k = 1, 2, 4, or 8

## Example

```
long m12(long x)
{
   return x*12;
}
```

## Converted to ASM by compiler:

```
leaq (%rdi,%rdi,2), %rax # t = x+2*x
salq $2, %rax # return t<<2</pre>
```

## **Some Arithmetic Operations**

#### Two Operand Instructions:

Computation		
Src,Dest	Dest = Dest + Src	
Src,Dest	Dest = Dest – Src	
Src,Dest	Dest = Dest * Src	
Src,Dest	Dest = Dest << Src	Also called shlq
Src,Dest	Dest = Dest >> Src	Arithmetic
Src,Dest	Dest = Dest >> Src	Logical
Src,Dest	Dest = Dest ^ Src	
Src,Dest	Dest = Dest & Src	
Src,Dest	Dest = Dest   Src	
	Src,Dest Src,Dest Src,Dest Src,Dest Src,Dest Src,Dest Src,Dest Src,Dest Src,Dest	Src,DestDest = Dest + SrcSrc,DestDest = Dest - SrcSrc,DestDest = Dest * SrcSrc,DestDest = Dest >> SrcSrc,DestDest = Dest >> SrcSrc,DestDest = Dest >> SrcSrc,DestDest = Dest ^ SrcSrc,DestDest = Dest & Src

- Watch out for argument order! Src,Dest
   (Warning: Intel docs use "op Dest,Src")
- No distinction between signed and unsigned int (why?)

## **Quiz Time!**

Check out:

https://canvas.cmu.edu/courses/5835

## **Some Arithmetic Operations**

#### One Operand Instructions

```
incq Dest Dest = Dest + 1

decq Dest Dest = Dest - 1

negq Dest Dest = -Dest

notq Dest Dest = \sim Dest
```

See book for more instructions

## **Arithmetic Expression Example**

```
long arith
(long x, long y, long z)
{
  long t1 = x+y;
  long t2 = z+t1;
  long t3 = x+4;
  long t4 = y * 48;
  long t5 = t3 + t4;
  long rval = t2 * t5;
  return rval;
}
```

```
arith:
  leaq (%rdi,%rsi), %rax
  addq %rdx, %rax
  leaq (%rsi,%rsi,2), %rdx
  salq $4, %rdx
  leaq 4(%rdi,%rdx), %rcx
  imulq %rcx, %rax
  ret
```

#### **Interesting Instructions**

- leaq: address computation
- **salq**: shift
- imulq: multiplication
  - But, only used once

# Understanding Arithmetic Expression Example

```
long arith
(long x, long y, long z)
{
  long t1 = x+y;
  long t2 = z+t1;
  long t3 = x+4;
  long t4 = y * 48;
  long t5 = t3 + t4;
  long rval = t2 * t5;
  return rval;
}
```

```
arith:
  leaq (%rdi,%rsi), %rax # t1
  addq %rdx, %rax # t2
  leaq (%rsi,%rsi,2), %rdx
  salq $4, %rdx # t4
  leaq 4(%rdi,%rdx), %rcx # t5
  imulq %rcx, %rax # rval
  ret
```

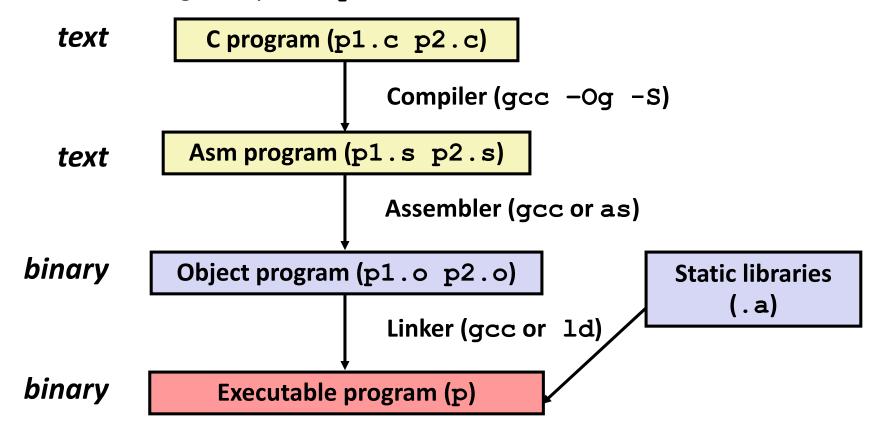
Register	Use(s)
%rdi	Argument x
%rsi	Argument <b>y</b>
%rdx	Argument <b>z</b> , <b>t4</b>
%rax	t1, t2, rval
%rcx	t5

# **Today: Machine Programming I: Basics**

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- C, assembly, machine code

## **Turning C into Object Code**

- Code in files p1.c p2.c
- Compile with command: gcc -Og p1.c p2.c -o p
  - Use basic optimizations (-Og) [New to recent versions of GCC]
  - Put resulting binary in file p



# **Compiling Into Assembly**

C Code (sum.c)

## **Generated x86-64 Assembly**

```
sumstore:
   pushq %rbx
   movq %rdx, %rbx
   call plus
   movq %rax, (%rbx)
   popq %rbx
   ret
```

Obtain (on shark machine) with command

Produces file sum.s

Warning: Will get very different results on non-Shark machines (Andrew Linux, Mac OS-X, ...) due to different versions of gcc and different compiler settings.

# What it really looks like

```
.globl sumstore
       .type sumstore, @function
sumstore:
.LFB35:
       .cfi startproc
       pushq %rbx
       .cfi def cfa offset 16
       .cfi offset 3, -16
       movq %rdx, %rbx
       call plus
       movq %rax, (%rbx)
       popq %rbx
       .cfi def cfa offset 8
       ret
       .cfi endproc
.LFE35:
       .size sumstore, .-sumstore
```

# What it really looks like

```
.glob1 sumstore
.type sumstore, @function
```

Things that look weird and are preceded by a "are generally directives.

```
sumstore:
.LFB35:
       .cfi startproc
       pushq %rbx
       .cfi def cfa offset 16
       .cfi offset 3, -16
       movq %rdx, %rbx
       call plus
       movq %rax, (%rbx)
       popq %rbx
       .cfi def cfa offset 8
       ret
       .cfi endproc
.LFE35:
       .size sumstore, .-sumstore
```

```
sumstore:
   pushq %rbx
   movq %rdx, %rbx
   call plus
   movq %rax, (%rbx)
   popq %rbx
   ret
```

## **Assembly Characteristics: Data Types**

- "Integer" data of 1, 2, 4, or 8 bytes
  - Data values
  - Addresses (untyped pointers)
- **■** Floating point data of 4, 8, or 10 bytes
- (SIMD vector data types of 8, 16, 32 or 64 bytes)
- Code: Byte sequences encoding series of instructions
- No aggregate types such as arrays or structures
  - Just contiguously allocated bytes in memory

## **Assembly Characteristics: Operations**

- Transfer data between memory and register
  - Load data from memory into register
  - Store register data into memory
- Perform arithmetic function on register or memory data
- Transfer control
  - Unconditional jumps to/from procedures
  - Conditional branches

# **Object Code**

#### Code for sumstore

#### 0x0400595: 0x53 0x48 0x89

0xd3

0xe8

0xf2

0xff

0xff

0xff

0x48

0x89

0x03

0x5b

0xc3

#### Assembler

- Translates .s into .o
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

#### Linker

- Resolves references between files
- Combines with static run-time libraries
  - E.g., code for malloc, printf
- Some libraries are dynamically linked
  - Linking occurs when program begins execution

Total of 14 bytes

Each instruction

1, 3, or 5 bytes

Starts at address

 $0 \times 0400595$ 

# **Machine Instruction Example**

0x40059e: 48 89 03

#### C Code

Store value t where designated by dest

## Assembly

- Move 8-byte value to memory
  - Quad words in x86-64 parlance
- Operands:

t: Register %rax

dest: Register %rbx

\*dest: Memory M[%rbx]

## Object Code

- 3-byte instruction
- Stored at address 0x40059e

## **Disassembling Object Code**

#### Disassembled

```
0000000000400595 <sumstore>:
  400595:
          53
                            push
                                   %rbx
  400596: 48 89 d3
                                   %rdx,%rbx
                           mov
  400599: e8 f2 ff ff ff
                           callq 400590 <plus>
  40059e: 48 89 03
                                   %rax, (%rbx)
                           mov
  4005a1:
          5b
                                   %rbx
                           pop
  4005a2: c3
                            reta
```

#### Disassembler

```
objdump -d sum
```

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can be run on either a .out (complete executable) or .o file

## **Alternate Disassembly**

#### Disassembled

```
Dump of assembler code for function sumstore:
    0x0000000000400595 <+0>: push      %rbx
    0x000000000400596 <+1>: mov      %rdx,%rbx
    0x000000000400599 <+4>: callq      0x400590 <plus>
    0x00000000040059e <+9>: mov      %rax,(%rbx)
    0x0000000004005a1 <+12>:pop      %rbx
    0x00000000004005a2 <+13>:retq
```

#### Within gdb Debugger

Disassemble procedure

```
gdb sum
disassemble sumstore
```

# **Alternate Disassembly**

## Object Code

## $0 \times 0400595$ : 0x530x480x890xd30xe8 0xf20xff 0xff 0xff0x480x890x03

0x5b

0xc3

#### **Disassembled**

```
Dump of assembler code for function sumstore:

0x00000000000400595 <+0>: push %rbx

0x0000000000400596 <+1>: mov %rdx,%rbx

0x0000000000400599 <+4>: callq 0x400590 <plus>
0x000000000040059e <+9>: mov %rax,(%rbx)

0x000000000004005a1 <+12>:pop %rbx

0x000000000004005a2 <+13>:retq
```

#### Within gdb Debugger

Disassemble procedure

gdb sum

disassemble sumstore

Examine the 14 bytes starting at sumstore

x/14xb sumstore

## What Can be Disassembled?

```
% objdump -d WINWORD.EXE
WINWORD.EXE: file format pei-i386
No symbols in "WINWORD.EXE".
Disassembly of section .text:
30001000 < text>:
30001000:
30001001:
               Reverse engineering forbidden by
30001003:
             Microsoft End User License Agreement
30001005:
3000100a:
```

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

## **Machine Programming I: Summary**

## History of Intel processors and architectures

Evolutionary design leads to many quirks and artifacts

## C, assembly, machine code

- New forms of visible state: program counter, registers, ...
- Compiler must transform statements, expressions, procedures into low-level instruction sequences

## Assembly Basics: Registers, operands, move

 The x86-64 move instructions cover wide range of data movement forms

#### Arithmetic

 C compiler will figure out different instruction combinations to carry out computation